CUDA-to-CSL remains a frontier-model capability

Full 28-kernel sweeps are solid bars; fixed 12-kernel probes are hatched.

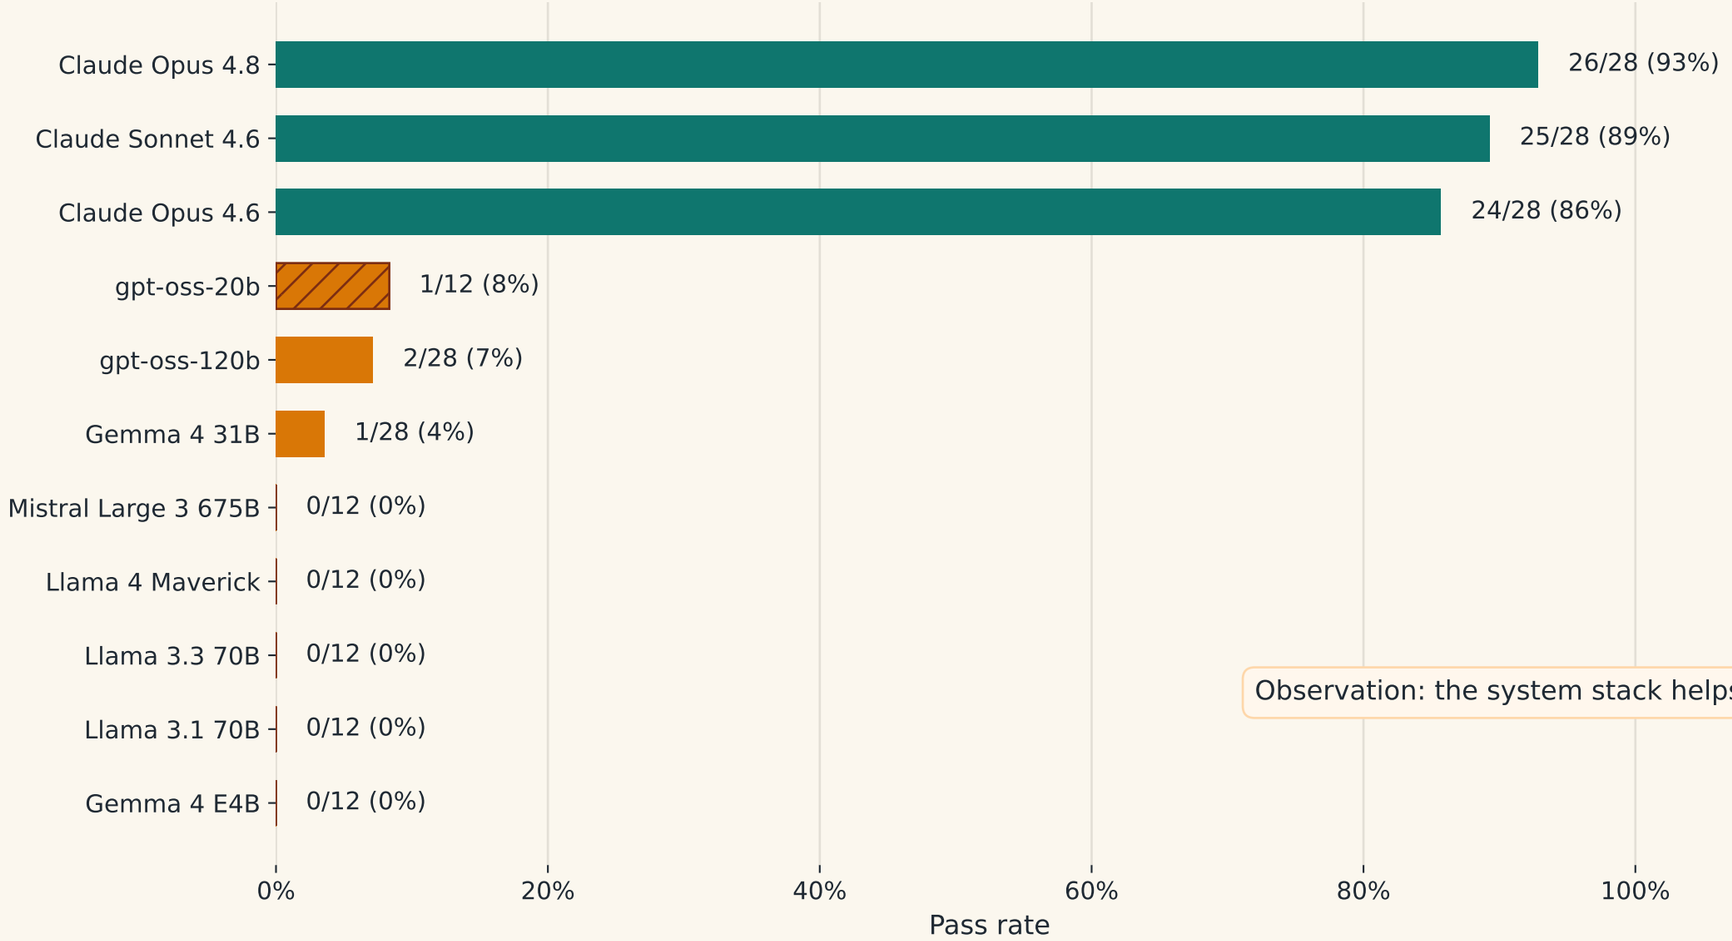

Observation: the system stack helps, but base CSL-generation skill is still the gate.

Sources: SWEEP_RESULTS_20260630.md and SESSION_PROGRESS_20260702.md. Rates with different denominators should not be over-compared.